

1.Run SQL connector

wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz

tar -xvf mysql-connector-java-8.0.25.tar.gz

cd mysql-connector-java-8.0.25/

sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/

```
login as: hadoop
Authenticating with public key "imported-openssh-key"

      _ _ _ _ _
     /   _   \   Amazon Linux 2 AMI
    / _ _   \
   /   _   \

https://aws.amazon.com/amazon-linux-2/

EEEEEEEEEEEEEEEEEEEE MMMMMMM          MMMMMMM RRRRRRRRRRRRRR
E::::::::::::::::::::E M::::::::M          M::::::::M R:::::::::R
EE::::::::EEEEEEEE::E M::::::::M          M::::::::M R:::RRRRR:::R
E:::E      EEEEE M::::::::M          M::::::::M RR:::R      R:::R
E:::E      EEEEE M::::::::M          M::::::::M R:::R      R:::R
E:::EEEEEEEEEE M:::M M:::M M:::M M:::M R:::RRRRR:::R
E::::::::::::E M:::M M:::M M:::M M:::M R:::RRRRR:::R
E:::EEEEEEEEEE M:::M M:::M M:::M M:::M R:::RRRRR:::R
E:::E      EEEEE M:::M M:::M M:::M R:::R      R:::R
EE:::EEEEEEEE::E M:::M          M:::M M:::M R:::R      R:::R
E::::::::::::E M:::M          M:::M RR:::R      R:::R
EEEEEEEEEEEEEEEEEEEE MMMMMMM          MMMMMMM RRRRRRR      RRRRRR

[hadoop@ip-172-31-35-158 ~]$ wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
--2023-01-22 13:58:45-- https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
Resolving de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)... 52.217.194.121, 54.231.193.65, 52.216.216.33, ...
Connecting to de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)|52.217.194.121|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 4079310 (3.9M) [application/x-gzip]
Saving to: 'mysql-connector-java-8.0.25.tar.gz'

100%[=====] 4,079,310 --.-K/s in 0.07s

2023-01-22 13:58:45 (57.4 MB/s) - 'mysql-connector-java-8.0.25.tar.gz' saved [4079310/4079310]

[hadoop@ip-172-31-35-158 ~]$ tar -xvf mysql-connector-java-8.0.25.tar.gz
mysql-connector-java-8.0.25/
mysql-connector-java-8.0.25/src/
mysql-connector-java-8.0.25/src/build/
```

```
mysql-connector-java-8.0.25/src/main/core-api/java/com/mysql/cj/conf/
mysql-connector-java-8.0.25/src/main/core-api/java/com/mysql/cj/exceptions/
mysql-connector-java-8.0.25/src/main/core-api/java/com/mysql/cj/interceptors/
mysql-connector-java-8.0.25/src/main/core-api/java/com/mysql/cj/log/
mysql-connector-java-8.0.25/src/main/core-api/java/com/mysql/cj/protocol/
mysql-connector-java-8.0.25/src/main/core-api/java/com/mysql/cj/result/
mysql-connector-java-8.0.25/src/main/core-api/java/com/mysql/cj/util/
mysql-connector-java-8.0.25/src/main/core-impl/
mysql-connector-java-8.0.25/src/main/core-impl/java/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/admin/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/conf/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/conf/url/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/log/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/protocol/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/protocol/result/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/result/
mysql-connector-java-8.0.25/src/main/core-impl/java/com/mysql/cj/sasl/
mysql-connector-java-8.0.25/src/main/doc/
mysql-connector-java-8.0.25/src/main/protocol-impl/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/mysql/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/mysql/cj/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/mysql/cj/protocol/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/mysql/cj/protocol/a/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/mysql/cj/protocol/a/authentication/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/mysql/cj/protocol/a/result/
mysql-connector-java-8.0.25/src/main/protocol-impl/java/com/mysql/cj/protocol/x/
mysql-connector-java-8.0.25/src/main/resources/
mysql-connector-java-8.0.25/src/main/resources/com/
mysql-connector-java-8.0.25/src/main/resources/com/mysql/
mysql-connector-java-8.0.25/src/main/resources/com/mysql/cj/
mysql-connector-java-8.0.25/src/main/resources/com/mysql/cj/configurations/
mysql-connector-java-8.0.25/src/main/resources/com/mysql/cj/util/
mysql-connector-java-8.0.25/src/main/user-api/
mysql-connector-java-8.0.25/src/main/user-api/java/
mysql-connector-java-8.0.25/src/main/user-api/java/com/
mysql-connector-java-8.0.25/src/main/user-api/java/com/mysql/
```

```
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/CompressionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/DevApiBaseTestCase.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/Ipv6SupportTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/MetadataTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/ResultTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SchemaTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SecureSessionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SessionFailoverTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SessionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableDeleteTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableInsertTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableSelectTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableUpdateTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TransactionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/package-info.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/InternalXBaseTestCase.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/MysqlxSessionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolAsyncTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolAuthTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/package-info.java
[hadoop@ip-172-31-35-158 ~]$ cd mysql-connector-java-8.0.25/
[hadoop@ip-172-31-35-158 mysql-connector-java-8.0.25]$ sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
cp: cannot create regular file '/usr/lib/sqoop/lib/': No such file or directory
[hadoop@ip-172-31-35-158 mysql-connector-java-8.0.25]$ sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
[hadoop@ip-172-31-35-158 mysql-connector-java-8.0.25]$
```

2. creation of HBase table and Sqoop command to ingest data from RDS into the HBase table: -

sqoop import \

--connect "jdbc:mysql://mysqlassignment1.cuyf7nwfv5s.us-east-1.rds.amazonaws.com:3306/demo" \

--username srbhrungta --password srbhrungta \

--table nyc_yellow_taxi \

--columns

"VendorID,tpep_pickup_datetime,tpep_dropoff_datetime,Passenger_count,Trip_distance,RateCodeID,Store_and_fwd_flag,PULocationID,DOLocationID,Payment_type,\

Fare_amount,Extra,MTA_tax,Tip_amount,Tolls_amount,Improvement_surcharge,Total_amount" \

--hbase-create-table \

--hbase-table nyc_yellow_taxi \

--column-family Trip_details \

--hbase-row-key VendorID,tpep_pickup_datetime,tpep_dropoff_datetime \

-m 1


```

SLF4J: Found binding in [jar:file:/usr/share/aws/redshift/jdbc/redshift-jdbc42-1.2.10.1041.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/lib/java/lib/loq4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
23/01/19 16:46:22 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
23/01/19 16:46:22 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
23/01/19 16:46:22 INFO tool.CodeGenerator: Beginning code generation
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.
23/01/19 16:46:24 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'nyc_yellow_taxi' AS t LIMIT 1
23/01/19 16:46:24 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'nyc_yellow_taxi' AS t LIMIT 1
23/01/19 16:46:24 INFO orm.CompilationManager: HADOOP MAPREDUCER is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-root/compile/1d4e93136935e7ad7b5932ad71a513e5/nyc_yellow_taxi.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
23/01/19 16:46:30 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-root/compile/1d4e93136935e7ad7b5932ad71a513e5/nyc_yellow_taxi.jar
23/01/19 16:46:30 WARN manager.MySQLManager: It looks like you are importing from mysql.
23/01/19 16:46:30 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
23/01/19 16:46:30 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
23/01/19 16:46:30 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
23/01/19 16:46:30 INFO mapreduce.ImportJobBase: Beginning import of nyc_yellow_taxi
23/01/19 16:46:30 INFO Configuration: deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
23/01/19 16:46:30 INFO Configuration: deprecation: mapred.map.task is deprecated. Instead, use mapreduce.job.maps
23/01/19 16:46:30 INFO mapreduce.HBaseImportJob: Creating missing HBase table nyc_yellow_taxi
23/01/19 16:46:35 WARN mapreduce.TableMapReduceUtil: The addDependencyJar(Configuration, Class<?>...) method has been deprecated since it is easy to use incorrectly. W
get users should rely on addDependencyJar(Job) instead. See HBASE-8216 for more details.
23/01/19 16:46:36 INFO client.RMProxy: Connecting to ResourceManager at ip-172-31-45-200.ec2.internal:172.31.49.200:8032
23/01/19 16:46:44 INFO dh.DHInputFormat: Using read committed transaction isolation.
23/01/19 16:46:44 INFO mapreduce.JobSubmitter: number of splits=1
23/01/19 16:46:44 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1674146772261_0001
23/01/19 16:46:45 INFO impl.YarnClientImpl: Submitted application application_1674146772261_0001
23/01/19 16:46:45 INFO mapreduce.Job: The url to track the job: http://ip-172-31-45-200.ec2.internal:8088/proxy/application_1674146772261_0001/
23/01/19 16:46:45 INFO mapreduce.Job: Running job: job_1674146772261_0001
23/01/19 16:46:45 INFO mapreduce.Job: Job job_1674146772261_0001 running in uber mode = false
23/01/19 16:46:45 INFO mapreduce.Job: map 0% reduce 0%
23/01/19 17:22:04 INFO mapreduce.Job: map 100% reduce 0%
23/01/19 17:22:05 INFO mapreduce.Job: Job job_1674146772261_0001 completed successfully
23/01/19 17:22:06 INFO mapreduce.Job: Counters: 39
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=224989
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=87
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=1
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Job Counters
    Launched map tasks=1
    Other local map tasks=1
    Total time spent by all maps in occupied slots (ms)=10092948
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=3122954
    Total reduce-millisecods taken by all map tasks=225244
    Total megabyte-millisecods taken by all map tasks=3230043616
  Map-Reduce Framework
    Map input records=18880595
    Map output records=18880595
    Input split bytes=87
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=19123
    CPU time spent (ms)=942240
    Physical memory (bytes) snapshot=61664696
    Virtual memory (bytes) snapshot=837897852
    Total committed heap usage (bytask) (bytes)=67414688
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=0
23/01/19 17:22:06 INFO mapreduce.ImportJobBase: Transferred 0 bytes in 2.1303122 seconds (0 bytes/sec)
23/01/19 17:22:06 INFO mapreduce.ImportJobBase: Retrieved 18880595 records.
[ec2-user@ip-172-31-45-200 ~]$

```

```

23/01/19 16:46:45 INFO mapreduce.Job: Running job: job_1674146772261_0001
23/01/19 16:46:45 INFO mapreduce.Job: Job job_1674146772261_0001 RUNNING in uber mode = false
23/01/19 16:46:45 INFO mapreduce.Job: map 0% reduce 0%
23/01/19 17:22:04 INFO mapreduce.Job: map 100% reduce 0%
23/01/19 17:22:05 INFO mapreduce.Job: Job job_1674146772261_0001 completed successfully
23/01/19 17:22:06 INFO mapreduce.Job: Counters: 39
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=224989
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=87
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=1
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Job Counters
    Launched map tasks=1
    Other local map tasks=1
    Total time spent by all maps in occupied slots (ms)=10092948
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=3122954
    Total reduce-millisecods taken by all map tasks=225244
    Total megabyte-millisecods taken by all map tasks=3230043616
  Map-Reduce Framework
    Map input records=18880595
    Map output records=18880595
    Input split bytes=87
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=19123
    CPU time spent (ms)=942240
    Physical memory (bytes) snapshot=61664696
    Virtual memory (bytes) snapshot=837897852
    Total committed heap usage (bytask) (bytes)=67414688
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=0
23/01/19 17:22:06 INFO mapreduce.ImportJobBase: Transferred 0 bytes in 2.1303122 seconds (0 bytes/sec)
23/01/19 17:22:06 INFO mapreduce.ImportJobBase: Retrieved 18880595 records.
[ec2-user@ip-172-31-45-200 ~]$

```